

# Graph Neural Networks II

Jan Tönshoff, Timo Gervens, Prof. Dr. Martin Grohe

RWTH Aachen

*toenshoff@informatik.rwth-aachen.de*

*gervens@informatik.rwth-aachen.de*

November 30, 2022

# Message Passing

Input:

- Graph  $G = (V, E)$
- Node Features  $X_V : V \rightarrow \mathbb{R}^d$
- Edge Features  $X_E : E \rightarrow \mathbb{R}^{d'}$  (optional)

# Message Passing

Input:

- Graph  $G = (V, E)$
- Node Features  $X_V : V \rightarrow \mathbb{R}^d$
- Edge Features  $X_E : E \rightarrow \mathbb{R}^{d'}$  (optional)

Idea of Message Passing:

- Initial embedding  $h^0(v) = X_V(v)$ .
- Update latent vertex embedding  $h^\ell(v) \in \mathbb{R}^{d^\ell}$  in each layer  $\ell$ .

# Message Passing

Input:

- Graph  $G = (V, E)$
- Node Features  $X_V : V \rightarrow \mathbb{R}^d$
- Edge Features  $X_E : E \rightarrow \mathbb{R}^{d'}$  (optional)

Idea of Message Passing:

- Initial embedding  $h^0(v) = X_V(v)$ .
- Update latent vertex embedding  $h^\ell(v) \in \mathbb{R}^{d^\ell}$  in each layer  $\ell$ .
- The update depends on the states of the neighbors.
- Use shared local neural network for the update.

# Message Passing

Elements of a GNN Layer [Wu et al., 2020]:

- Trainable functions  $\mathbf{M}^\ell$ ,  $\mathbf{U}^\ell$
- Aggregation function  $\oplus$  (Sum, Mean, Max, ...)

## Message Passing

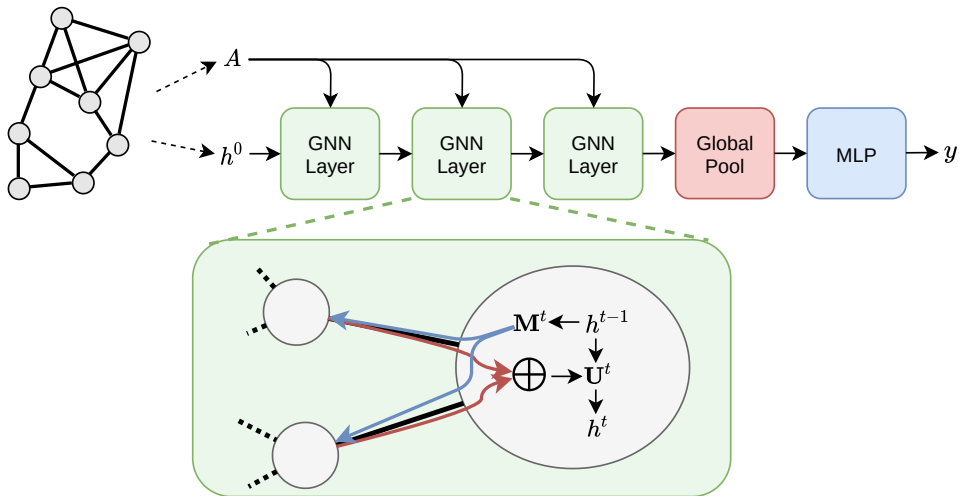
Elements of a GNN Layer [Wu et al., 2020]:

- Trainable functions  $\mathbf{M}^\ell$ ,  $\mathbf{U}^\ell$
- Aggregation function  $\bigoplus$  (Sum, Mean, Max, ...)

Function computed by the layer:

$$h^\ell(v) = \mathbf{U}^\ell \left( h^{\ell-1}(v), \bigoplus_{u \in \mathcal{N}(v)} \mathbf{M}^\ell \left( h^{\ell-1}(u), X_E(vu) \right) \right)$$

# Graph Neural Networks



## Implementing Message Passing

How do you actually implement Message Passing in PyTorch?



## Implementing Message Passing

How do you actually implement Message Passing in PyTorch?

Simple Solution: Dense Matrix Multiplication

- Easy to implement
- Highly optimized for small graphs

## Implementing Message Passing

How do you actually implement Message Passing in PyTorch?

Simple Solution: Dense Matrix Multiplication

- Easy to implement
- Highly optimized for small graphs
- Does not scale to large graphs
- Requires padding

## Implementing Message Passing

How do you actually implement Message Passing in PyTorch?

Simple Solution: Dense Matrix Multiplication

- Easy to implement
- Highly optimized for small graphs
- Does not scale to large graphs
- Requires padding
- Only linear aggregations
- No edge features

# Implementing Message Passing

How do you actually implement Message Passing in PyTorch?

Simple Solution: Dense Matrix Multiplication

- Easy to implement
- Highly optimized for small graphs
- Does not scale to large graphs
- Requires padding
- Only linear aggregations
- No edge features

Better Solution: **Scatter Operations**

## PyTorch Scatter

Open Source Torch Extension<sup>1</sup>.

Compatible with CUDA and Backpropagation.

---

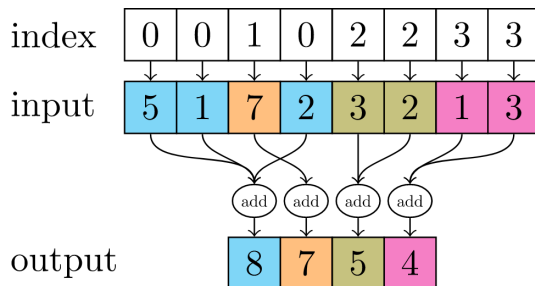
<sup>1</sup>[https://github.com/rusty1s/pytorch\\_scatter](https://github.com/rusty1s/pytorch_scatter)

## PyTorch Scatter

Open Source Torch Extension<sup>1</sup>.

Compatible with CUDA and Backpropagation.

Enables aggregation of input with index list:



<sup>1</sup>[https://github.com/rusty1s/pytorch\\_scatter](https://github.com/rusty1s/pytorch_scatter)

# PyTorch Scatter Example

$$x = \begin{pmatrix} 0 & 2 \\ 1 & 0 \\ 1 & 1 \\ 3 & 5 \\ 0 & 1 \end{pmatrix}$$

$$\text{idx} = [0, 1, 0, 2, 2]$$

## PyTorch Scatter Example

$$x = \begin{pmatrix} 0 & 2 \\ 1 & 0 \\ 1 & 1 \\ 3 & 5 \\ 0 & 1 \end{pmatrix}$$

`scatter_sum(x, idx, dim=0) =`

$$\text{idx} = [0, 1, 0, 2, 2]$$



## PyTorch Scatter Example

$$x = \begin{pmatrix} 0 & 2 \\ 1 & 0 \\ 1 & 1 \\ 3 & 5 \\ 0 & 1 \end{pmatrix}$$

$$\text{scatter\_sum}(x, \text{idx}, \text{dim}=0) = \begin{pmatrix} 1 & 3 \\ 1 & 0 \\ 3 & 6 \end{pmatrix}$$

$$\text{idx} = [0, 1, 0, 2, 2]$$

## PyTorch Scatter Example

$$x = \begin{pmatrix} 0 & 2 \\ 1 & 0 \\ 1 & 1 \\ 3 & 5 \\ 0 & 1 \end{pmatrix}$$

$$\text{idx} = [0, 1, 0, 2, 2]$$

$$\text{scatter\_sum}(x, \text{idx}, \text{dim}=0) = \begin{pmatrix} 1 & 3 \\ 1 & 0 \\ 3 & 6 \end{pmatrix}$$

$$\text{scatter\_max}(x, \text{idx}, \text{dim}=0)[0] =$$

## PyTorch Scatter Example

$$x = \begin{pmatrix} 0 & 2 \\ 1 & 0 \\ 1 & 1 \\ 3 & 5 \\ 0 & 1 \end{pmatrix}$$

$$idx = [0, 1, 0, 2, 2]$$

$$\text{scatter\_sum}(x, idx, \text{dim}=0) = \begin{pmatrix} 1 & 3 \\ 1 & 0 \\ 3 & 6 \end{pmatrix}$$

$$\text{scatter\_max}(x, idx, \text{dim}=0)[0] = \begin{pmatrix} 1 & 2 \\ 1 & 0 \\ 3 & 5 \end{pmatrix}$$

# Sparse Graph Representation

Sparse representation of  $G = (V, E)$ :

# Sparse Graph Representation

Sparse representation of  $G = (V, E)$ :

1. *Directed* edge list  $\text{idx}_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$

# Sparse Graph Representation

Sparse representation of  $G = (V, E)$ :

1. *Directed* edge list  $\text{id}_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$
2. Node feature matrix  $X_V \in \mathbb{R}^{|V| \times d}$
3. Optional edge feature matrix  $X_E \in \mathbb{R}^{2|E| \times d'}$

## Sparse Graph Representation

Sparse representation of  $G = (V, E)$ :

1. *Directed* edge list  $\text{idx}_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$
2. Node feature matrix  $X_V \in \mathbb{R}^{|V| \times d}$
3. Optional edge feature matrix  $X_E \in \mathbb{R}^{2|E| \times d'}$
4. Graph label  $y \in \mathbb{R}^c$

my guess  
j'th column gives you source  
and target of j'th edge  
(u)  
(v)

## Message Passing with Scatter

Given:  $H^{(\ell-1)} \in \mathbb{R}^{|V| \times d_h}$ ,  $X_E \in \mathbb{R}^{2|E| \times d'}$ ,  $\text{idx}_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$

Compute:



## Message Passing with Scatter

Given:  $H^{(\ell-1)} \in \mathbb{R}^{|V| \times d_h}$ ,  $X_E \in \mathbb{R}^{2|E| \times d'}$ ,  $\text{idx}_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$

Compute:

$$Y^{(\ell)} = \mathbf{M}(H^{(\ell-1)}[\text{idx}_E[0]], X_E)$$

## Message Passing with Scatter

Given:  $H^{(\ell-1)} \in \mathbb{R}^{|V| \times d_h}$ ,  $X_E \in \mathbb{R}^{2|E| \times d'}$ ,  $\text{id}x_E \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$

Compute:

$$Y^{(\ell)} = \mathbf{M}(H^{(\ell-1)}[\text{id}x_E[0]], X_E)$$

$$Z^{(\ell)} = \text{scatter\_sum}(Y^{(\ell)}, \text{id}x_E[1], \text{dim}=0)$$

## Message Passing with Scatter

Given:  $H^{(\ell-1)} \in \mathbb{R}^{|V| \times d_h}$ ,  $X_E \in \mathbb{R}^{2|E| \times d'}$ ,  $\text{idxE} \in \{0, \dots, |V| - 1\}^{2 \times 2|E|}$

Compute:

$$Y^{(\ell)} = \mathbf{M}(H^{(\ell-1)}[\text{idxE}[0]], X_E)$$

$H[\text{idxs}[0]] - (2|E|, d_h)$   
 $H\_Xe\_ - (2|E|, d_h + d')$

$$Z^{(\ell)} = \text{scatter\_sum}(Y^{(\ell)}, \text{idxE}[1], \text{dim}=0)$$

$$H^{(\ell)} = \mathbf{U}(H^{(\ell-1)}, Z^{(\ell)})$$

## Batching Sparse Graphs

Given: Batch of training graphs  $(\text{idx}_{E_1}, X_{V_1}, X_{E_1}, y_1), \dots, (\text{idx}_{E_b}, X_{V_b}, X_{E_b}, y_b)$



## Batching Sparse Graphs

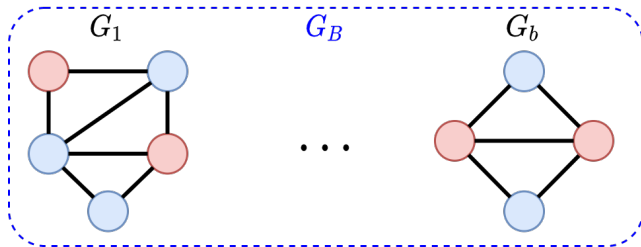
Given: Batch of training graphs  $(\text{idx}_{E_1}, X_{V_1}, X_{E_1}, y_1), \dots, (\text{idx}_{E_b}, X_{V_b}, X_{E_b}, y_b)$



How do we combine these into one input for parallel processing?

## Batching Sparse Graphs

Given: Batch of training graphs  $(\text{idx}_{E_1}, X_{V_1}, X_{E_1}, y_1), \dots, (\text{idx}_{E_b}, X_{V_b}, X_{E_b}, y_b)$



How do we combine these into one input for parallel processing?

Compute disjoint union  $G_B = (V_B, E_B)$ :  $V_B = \bigcup_{i=1}^b V_i$ ,  $E_B = \bigcup_{i=1}^b E_i$

## Batching Sparse Graphs

Representation of the batched graph  $G_B$ :

1.  $\text{idx}_E \in \{0, \dots, |V_B| - 1\}^{2 \times 2|E_B|}$
2.  $X_{V_B} \in \mathbb{R}^{|V_B| \times d}$
3.  $X_{E_B} \in \mathbb{R}^{|E_B| \times d'}$
4.  $y_B \in \mathbb{R}^{B \times c}$
5.  $\text{batch\_idx} \in \{0, \dots, b - 1\}^{|V_B|}$  (remember original graph for each node)

## Batching Sparse Graphs

Representation of the batched graph  $G_B$ :

1.  $\text{idx}_E \in \{0, \dots, |V_B| - 1\}^{2 \times 2|E_B|}$
2.  $X_{V_B} \in \mathbb{R}^{|V_B| \times d}$
3.  $X_{E_B} \in \mathbb{R}^{|E_B| \times d'}$
4.  $y_B \in \mathbb{R}^{B \times c}$
5.  $\text{batch\_idx} \in \{0, \dots, b - 1\}^{|V_B|}$  (remember original graph for each node)

The indices in  $\text{idx}_{E_B}$  are re-enumerated.



## Batching Sparse Graphs

Representation of the batched graph  $G_B$ :

1.  $\text{idx}_E \in \{0, \dots, |V_B| - 1\}^{2 \times 2|E_B|}$
2.  $X_{V_B} \in \mathbb{R}^{|V_B| \times d}$
3.  $X_{E_B} \in \mathbb{R}^{|E_B| \times d'}$
4.  $y_B \in \mathbb{R}^{B \times c}$
5.  $\text{batch\_idx} \in \{0, \dots, b - 1\}^{|V_B|}$  (remember original graph for each node)

The indices in  $\text{idx}_{E_B}$  are re-enumerated.

Pooling becomes a scatter operation with  $\text{batch\_idx}$ .

## Batching Sparse Graphs

Representation of the batched graph  $G_B$ :

1.  $\text{idx}_E \in \{0, \dots, |V_B| - 1\}^{2 \times 2|E_B|}$
2.  $X_{V_B} \in \mathbb{R}^{|V_B| \times d}$
3.  $X_{E_B} \in \mathbb{R}^{|E_B| \times d'}$
4.  $y_B \in \mathbb{R}^{B \times c}$
5.  $\text{batch\_idx} \in \{0, \dots, b - 1\}^{|V_B|}$  (remember original graph for each node)

The indices in  $\text{idx}_{E_B}$  are re-enumerated.

Pooling becomes a scatter operation with  $\text{batch\_idx}$ .

## Batching Sparse Graphs in PyTorch

Write custom collation function in Python:

- Input:  $[(\text{idx}_{E_1}, X_{V_1}, X_{E_1}, y_1), \dots, (\text{idx}_{E_b}, X_{V_b}, X_{E_b}, y_b)]$
- Output:  $(\text{idx}_{E_B}, X_{V_B}, X_{E_B}, y_B, \text{batch\_idx})$
- Pass function to `DataLoader` as `collate_fn` parameter.

## GNNs in Practice

Typical Hyperparameters:

- 3-10 layers
- Latent dimension  $50 \leq d \leq 500$

## GNNs in Practice

Typical Hyperparameters:

- 3-10 layers
- Latent dimension  $50 \leq d \leq 500$

Most GNNs are relatively shallow and small (compared to CNNs, Transformers, etc.)

## GNN Components

Update function **U**:

- Usually MLP with one or two layers
- If **U** has two layer, the second layer is usually linear

# GNN Components

Update function **U**:

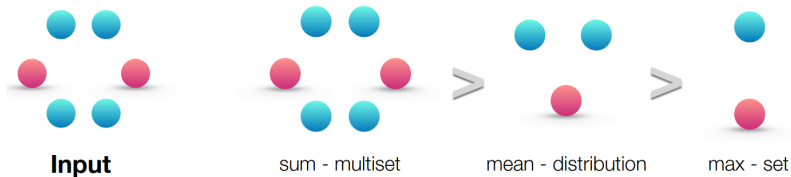
- Usually MLP with one or two layers
- If **U** has two layer, the second layer is usually linear

Message function **M**:

- If no edge features are given, simply use identity (like the GCN)
- If edge features are given, **M** should be non-linear.

# Aggregation

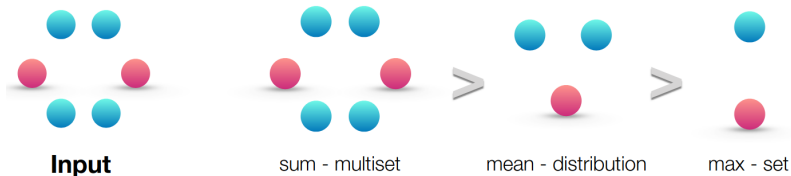
Expressiveness of aggregation functions [Xu et al., 2019]:





# Aggregation

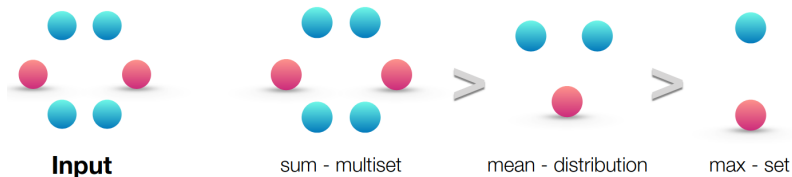
Expressiveness of aggregation functions [Xu et al., 2019]:



- SUM carries the most information but is not very robust.

## Aggregation

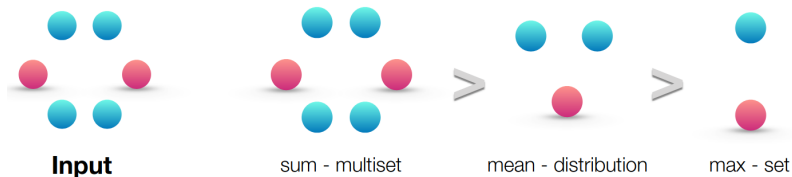
Expressiveness of aggregation functions [Xu et al., 2019]:



- SUM carries the most information but is not very robust.
- MEAN is more robust but you may need feature engineering.

# Aggregation

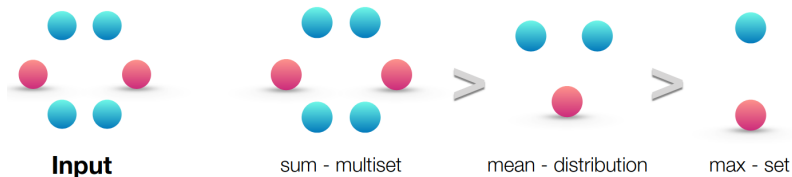
Expressiveness of aggregation functions [Xu et al., 2019]:



- SUM carries the most information but is not very robust.
- MEAN is more robust but you may need feature engineering.
- MAX carries the least information but can handle extreme ranges of degrees.

## Aggregation

Expressiveness of aggregation functions [Xu et al., 2019]:



- SUM carries the most information but is not very robust.
- MEAN is more robust but you may need feature engineering.
- MAX carries the least information but can handle extreme ranges of degrees.

More advanced: Predict edge weights with attention. [Veličković et al., 2018]

## Special GNNs

Recurrent GNNs:

- Choose **U** as GRU-Cell and reuse it in each message pass.
- Number of message passes not fixed.
- Application: Learning Heuristics for NP-hard problems.

## Special GNNs

### Recurrent GNNs:

- Choose **U** as GRU-Cell and reuse it in each message pass.
- Number of message passes not fixed.
- Application: Learning Heuristics for NP-hard problems.

### Heterogeneous GNNs:

- Heterogeneous graphs contain nodes of different types.
- Train different functions **U**, **M** for each type of node.

## Standard Tricks

Most of the standard tricks from Deep Learning also work for GNNs:

- Batch/Layer Normalization
- Dropout

## Standard Tricks

Most of the standard tricks from Deep Learning also work for GNNs:

- Batch/Layer Normalization
- Dropout
- Use Residual Connections:

$$h^\ell(v) = h^{\ell-1}(v) + \mathbf{U}^\ell \left( h^{\ell-1}(v), \bigoplus_{u \in \mathcal{N}(v)} \mathbf{M}^\ell \left( h^{\ell-1}(u), X_E(vu) \right) \right)$$



## Virtual Nodes

Perform global pooling after each GNN layer and pass result back to nodes:

$$h^\ell(G) = \mathbf{V}^\ell \left( \sum_{v \in V} h^\ell(v) \right)$$
$$\tilde{h}^\ell(v) = h^\ell(v) + h^\ell(G)$$

Here,  $\mathbf{V}^\ell$  is a trainable MLP.

Virtual Nodes enable global information exchange after each layer.

## Local Tasks on Massive Graphs

Dataset: One huge graph with billions of vertices.

Task: Node Classification with a GNN.

## Local Tasks on Massive Graphs

Dataset: One huge graph with billions of vertices.

Task: Node Classification with a GNN.

Solution: Sample subgraph rooted in  $v$ . Only process this subgraph to predict  $y(v)$ .

## Local Tasks on Massive Graphs

Dataset: One huge graph with billions of vertices.

Task: Node Classification with a GNN.

Solution: Sample subgraph rooted in  $v$ . Only process this subgraph to predict  $y(v)$ .

Sampling Algorithms:

- GraphSAGE [Hamilton et al., 2017]
- HGSampler [Hu et al., 2020]

## Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

## Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

Extensions:

- Higher order GNNs (equivalent to higher order WL-Test) [Morris et al., 2019]

## Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

Extensions:

- Higher order GNNs (equivalent to higher order WL-Test) [Morris et al., 2019]
- Random node features [Sato et al., 2020]

# Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

Extensions:

- Higher order GNNs (equivalent to higher order WL-Test) [Morris et al., 2019]
- Random node features [Sato et al., 2020]
- Pre-Compute substructures of interest [Bouritsas et al., 2020]



# Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

Extensions:

- Higher order GNNs (equivalent to higher order WL-Test) [Morris et al., 2019]
- Random node features [Sato et al., 2020]
- Pre-Compute substructures of interest [Bouritsas et al., 2020]

Is this even the correct measure of “power”?

## Overcoming WL

Standard GNNs are at most as powerful as the (1-dimensional) Weisfeiler-Leman-Test.

Extensions:

- Higher order GNNs (equivalent to higher order WL-Test) [Morris et al., 2019]
- Random node features [Sato et al., 2020]
- Pre-Compute substructures of interest [Bouritsas et al., 2020]

Is this even the correct measure of “power”?

This is an active field of research.

## Exercise 3

Implement a more general GNN layer in PyTorch:

- Implement custom collation function.

## Exercise 3

Implement a more general GNN layer in PyTorch:

- Implement custom collation function.
- Implement GNN layer with scatter operations (SUM, MEAN and MAX).

## Exercise 3

Implement a more general GNN layer in PyTorch:

- Implement custom collation function.
- Implement GNN layer with scatter operations (SUM, MEAN and MAX).
- Implement global pooling with scatter operations.

## Exercise 3

Implement a more general GNN layer in PyTorch:

- Implement custom collation function.
- Implement GNN layer with scatter operations (SUM, MEAN and MAX).
- Implement global pooling with scatter operations.
- Implement a Virtual Node layer.

## Exercise 3

Implement a more general GNN layer in PyTorch:

- Implement custom collation function.
- Implement GNN layer with scatter operations (SUM, MEAN and MAX).
- Implement global pooling with scatter operations.
- Implement a Virtual Node layer.
- Evaluate your implementation on the ZINC dataset.

## References I

- Zonghan Wu, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and S Yu Philip. A comprehensive survey on graph neural networks. *IEEE transactions on neural networks and learning systems*, 2020.
- Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? In *Proceedings of the Seventh International Conference on Learning Representations (ICLR)*, 2019.
- Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. 2018.
- Will Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. *Advances in neural information processing systems*, 30, 2017.
- Ziniu Hu, Yuxiao Dong, Kuansan Wang, and Yizhou Sun. Heterogeneous graph transformer. In *Proceedings of The Web Conference 2020*, pages 2704–2710, 2020.



## References II

- Christopher Morris, Martin Ritzert, Matthias Fey, William L. Hamilton, Jan Eric Lenssen, Gaurav Rattan, and Martin Grohe. Weisfeiler and Leman go neural: Higher-order graph neural networks. In *Proceedings of the Thirty-Third AAAI Conference on Artificial Intelligence (AAAI)*, pages 4602–4609, 2019.
- Ryoma Sato, Makoto Yamada, and Hisashi Kashima. Random features strengthen graph neural networks. *CoRR*, abs/2002.03155, 2020.
- Giorgos Bouritsas, Fabrizio Frasca, Stefanos Zafeiriou, and Michael M Bronstein. Improving graph neural network expressivity via subgraph isomorphism counting. *arXiv preprint arXiv:2006.09252*, 2020.